



## ML Lab Week 14: CNN Image Classification

**Name:**Najmus Seher

**SRN:**PES2UG23CS359

**Section:**F

**Course:** Machine Learning

## 1. Introduction

This lab builds, trains, and evaluates a Convolutional Neural Network (CNN) to classify images of hands showing Rock, Paper, or Scissors. The objective is to prepare image data, design a small CNN, train it on the provided dataset, and measure its performance on unseen test images.

## 2. Model Architecture

### Convolutional feature extractor (conv\_block):

- The network uses three convolutional blocks. Each block contains:
  1. Conv2d with kernel size = 3, padding = 1, then ReLU, then MaxPool2d(2).
  2. Channel progression:
    - Block 1: input channels = 3 -> output channels = 16
    - Block 2: 16 -> 32
    - Block 3: 32 -> 64
- Because each MaxPool2d halves spatial dimensions, an input image resized to 128×128 becomes 64×64 after the first pool, 32 × 32 after the second, and 16×16 after the third.

### Fully-connected classifier (fc):

- After the conv blocks the feature map shape is (batch\_size, 64, 16, 16).
- The classifier is:
  1. Flatten()
  2. Linear( $64 * 16 * 16 \rightarrow 256$ )
  3. ReLU()
  4. Dropout( $p=0.3$ )
  5. Linear( $256 \rightarrow 3$ ) — final logits for the three classes (rock, paper, scissors).

**Design rationale:** compact architecture keeps model small (fast to train) while providing progressive feature abstraction (increasing channels + downsampling). ReLU activation and dropout chosen for nonlinearity and basic regularization.

## 3. Training and Performance

### Data processing & loader details

- Image transforms:

- Resize all images to  $128 \times 128$
- `ToTensor()`
- Normalize with mean = 0.5, std = 0.5 (applied per channel)
- Dataset loading:
  - Used `torchvision.datasets.ImageFolder(DATA_DIR, transform=transform)`
  - Split: 80% training, 20% test using `random_split`
  - Batch size = 32
  - `train_loader` shuffled, `test_loader` not shuffled

## **Training hyperparameters**

- Optimizer: Adam
- Loss function: CrossEntropyLoss
- Learning rate: 0.001
- Epochs: 10
- Batch size: 32
- Device: `torch.device("cuda" if torch.cuda.is_available() else "cpu")`-training uses GPU if available, otherwise CPU.

In my case I used the cpu

## **Final test accuracy**

Test Accuracy: 98.17%

## **4. Conclusion and Analysis**

### **Overall Performance**

- The model performed **very well**, achieving a **98.17% test accuracy**, which shows that the CNN learned to distinguish Rock, Paper, and Scissors images almost perfectly.
- The high accuracy also suggests that the model generalized well to unseen test images.

### **Challenges Faced**

- One challenge was making sure the dataset paths and folder structure were correct so the images could load without errors.
- Another challenge was understanding how the CNN layers work together, especially calculating output sizes for the fully connected layer.
- Training also took some time, especially when running on CPU, so monitoring the progress was important.

### **Possible Improvements**

- I could add **data augmentation** (like rotations, flips, brightness changes) to make the model even more robust to different hand positions or lighting conditions.
- Another improvement would be using **transfer learning** (e.g., ResNet18), which could slightly boost accuracy and reduce training time.
- Increasing the number of epochs or tuning hyperparameters (like learning rate or batch size) could also help push the accuracy even higher.